

Type Inference and Type Checking for JavaScript Strict Mode

Micha Reiser

University of Applied Sciences Rapperswil
Supervised by Luc Bläser
ST2016 (May 24, 2016)

Abstract

The popularity of JavaScript dramatically increased over the last years. It evolved from the programming language of the web to a general purpose language that is also used for mobile-, desktop- and server-applications. However, the tooling support for program verification is sparse due to the dynamic nature of JavaScript that is hard to be covered using static analysis. The existing verification tools are either limited to simple bug patterns, are based on a super or subset of JavaScript, or only support outdated JavaScript versions. This work describes an algorithm for type inference and type checking of JavaScript code written in strict mode. The defined algorithm combines the Hindley-Milner Algorithm W with abstract interpretation. The type system used by the algorithm is unsound as the precision of the type inference degrades for very dynamic code, code that is mainly located in frameworks or libraries. The defined algorithm is implemented in *ES-Checker* and is compared to competitive type checkers. The results show that a precise type inference can be achieved for a majority of programs and provides a valuable feedback for programmers.

1 Introduction

The role of JavaScript dramatically changed over the last years. From an unpopular language used to add dynamic effects to web pages to a widely used language with a strong and growing community. It emerged from a browser only language to a general purpose language used to write web-, desktop-, mobile-, and server-applications. This shift is reflected in an increasing complexity and number of JavaScript projects. To tackle the higher complexity, a better tooling support is needed for effective development and refactoring.

JavaScript does not provide any static analysis for proving the soundness of a program. Type and nullability checking is only performed at runtime. Performing refactorings or adding new functionality is therefore a risky task, as the programmer has a very limited tooling for testing if changes have been applied correctly. Static type checking allows to

detect common errors like accessing non-declared variables, missing arguments, arguments in incorrect order or invoking a non function type, without the need to execute the program. It therefore is a valuable tool for providing a fast statement about the soundness of a program.

JavaScript is a dynamically typed language and therefore requires type inference for type checking. Type inference for JavaScript is a non-trivial task because of the very dynamic nature of JavaScript [9]. JavaScript has several features that makes static program analysis difficult. An explanation of the language features adding complexity to static analysis follows.

Dynamic Objects A JavaScript object is a mapping from a string key to a value. Adding new properties or removing existing once can be performed dynamically. A property name can either be a static or computed value. The later is used for dynamic object creation or modification, similar to code using reflection in statically typed languages.

Closures Functions have access to variables from their enclosing scope. Invoking a function requires that the function is evaluated in its declaration scope. Therefore the analysis needs to be context-sensitive.

Side effects Functions in JavaScript are not pure and therefore invoking a function can have side effects to arguments passed to the function or variables from the outer scope of the function declaration. These side effects can also affect the type of the involved variables, e.g. if the function assigns a value of a different type to an outer scope variable or adds a new property to a parameter. A precise analysis needs to reflect these side effects in the caller's context.

This Binding The object referenced by *this* depends on the function kind and its usage. Arrow functions capture the *this* of the enclosing context. The *this* inside of a function declarations or expressions depends on the usage. The *this* can explicitly be specified if the function is invoked using *call* or *apply*. Otherwise the binding of *this* is implicitly defined. If the function is called as a method of an object, then *this* is equal to the object, to which the

method belongs. If the function is not a member of an object, then *this* references the global object.

Host-Environment The preliminary host-environment of JavaScript applications is still the browser. However, it is also used for standalone applications or to add scripting functionality to other applications. In this case, NodeJS, Rhino or another JavaScript engine is used as host-environment. Each environment exposes different native objects and methods at runtime, e.g. the browser exposes the DOM-API. A type checker needs to specifically model these objects and methods as their JavaScript code is not existent.

Implementing a sound type checker for JavaScript risks to be over-restrictive and only allows a very limited subset of JavaScript programs or reports a large amount of false positives. An unsound type system has the disadvantage that it does not detect all errors but allows type checking of a far more complete set of JavaScript programs and therefore has a better chance to be applied in actual practice.

This work defines a sound algorithm that is capable of inferring the types and perform type checking for a majority of the written code. The precision degrades for very dynamic code, e.g. code that uses dynamic object creation or manipulation. The algorithm is unsound for these cases as the inferred types might be imprecise. It is expected that these features are mainly used in frameworks or libraries. This work suggests to substitute type inference for these edge cases by using type annotations. The presented analysis is limited to JavaScript code written in strict mode. Features prohibited in strict mode are not supported by the analysis.

The first section compares this work with existing tools used to verify or analyze JavaScript programs. Section 3 explains the benefits of restricting the analysis to strict mode and why it is believed that the code written in strict mode will be growing in the near future. The basics of the algorithm are explained in section 4, the results from the evaluation are shown in the preceding section and is followed by the conclusion.

2 Related Work

This work is related to linters, transpiled languages and other type checkers for JavaScript.

Linters like ESLint [5] are used to enforce a specific coding style across a project or to find errors using common bug patterns. Linters use simple static analysis techniques for bug identification. These analysis mostly are intra-procedural. This work focuses on errors deducible by type checking, based on a sophisticated inter-procedural analysis.

An alternative and quite popular approach for type checking programs executing in a JavaScript environment is by transpiling a source language to JavaScript. The source language either allows type inference [2, 12] or uses type annotations. Well-known examples are TypeScript [13] and Flow [4]. The downside of a transpiled language is the need for an additional build step that slows down the development cycle and is a potential source for errors. A developer using

a transpiled language needs to have a good understanding of the source language and JavaScript. This limits the amount of potential programmers or requires additional training. This work differs from transpilers as the focus is on type checking JavaScript and not transpiled languages.

TAJS [9] is a sound type analyzer and type checker for JavaScript. The used algorithm is context and path-sensitive. TAJIS uses abstract interpretation for type inference. The goal of TAJIS is a precise and sound type checker that supports the full JavaScript language. The current version only supports ECMAScript 3 and a limited set of ECMAScript 5. Compared to TAJIS, this work focuses on JavaScript programs using ECMAScript 6 and strict mode. Furthermore, the presented analysis is unsound to reduce the number of false positives.

Infernu [11] implements type inference and type checking for JavaScript. It uses the Hindley-Milner algorithm. The used type system only models a limited set of JavaScript. The subset is defined by the properties required by the Hindley-Milner algorithm. For instance, the unmodified Hindley-Milner algorithm requires that a variable has exactly one type in a program, the principal type. Therefore, Infernu disallows assigning values of different types to the same variable, contrary to the JavaScript specification. This work differs from Infernu as it extends the Hindley-Milner algorithm to support a wider set of JavaScript programs and reduce the number of false positives.

Odgaard describes in his master thesis a dynamic analysis for type inference [16]. The presented analysis uses code instrumentation to obtain the runtime values and derives the variable types from these. The inferred types are used to add JSDoc [10] type annotations. Compared to the approach presented by the work of Odgaard, this work uses static over dynamic analysis for type inference.

Tern [6] is an editor-independent JavaScript analyzer with type inference. It only provides an API for editors but does not perform type checking. Editors can use the API of Tern to query type-information, provide auto-completion, and jump-to-definition functionality. Tern uses abstract values and abstract interpretation for type inference. Tern can only infer types for functions with an actual invocation. Non-invoked functions are not analyzed. This work differs from Tern as the project focuses on type checking and not on providing an API for editors.

3 Benefits of Strict Mode

Strict mode has been introduced in ECMAScript 5. Strict mode allows to opt in to a restricted variant of JavaScript. Strict mode is not only a subset of JavaScript, it intentionally changes the semantics from normal code. It eliminates silent errors by throwing exceptions instead. It also prohibits some error-prone and hard-to-optimize syntaxes and semantics from earlier ECMAScript versions.

Strict mode can be explicitly enabled by adding the *"use strict"* directive before any other statement in a file or function. Using the directive in a file enables strict mode for the whole file, using it in a function enables strict mode for a specific function. Strict mode is enforced for scripts using

ECMAScript 6 modules [7, p. 10.2.1]. Therefore, it can be expected that newer code is using strict mode. The analysis only supports code written in strict mode to take advantages of the changed semantics. A description of the changed semantics with an effect to the analysis follows.

Prohibited With Statement The *with* statement is prohibited in strict mode [7, Annex C]. Inside the *with* statement object-properties can be accessed without the need to use member expressions. In the following example, the identifier *x* on line three either references the property *obj.x* or the variable *x* defined on line one.

```
1  var x = 17;
2  with (obj) {
3    x; // references obj.x or variable x
4  }
```

If the object *obj* has a property *x*, then the identifier references the property, otherwise it references the variable *x*. This behavior makes lexical scoping a non-trivial task [9]. The prohibition of the *with* statement allows static scoping.

Assignment to non-declared Variables An Assignment to a non-declared variable introduces a global variable in non-strict mode. Strict mode prohibits assignments to non-declared variables and throws an error instead. Variables can not be implicitly declared in strict mode. Therefore, accessing a not yet known variable is always an error.

4 Algorithm

The classical approach for type inference is the Hindley-Milner algorithm [14]. The Hindley-Milner algorithm infers the principal type for every variable in a program. This is sufficient for languages restricting that the type of a variable does not change over its lifespan. In contrary, JavaScript has no such restriction allowing values of different types to be assigned to the same variable. For instance, a common JavaScript pattern is to first declare the variables and defer their initialization. The type of the variable after its declaration is *void*. The initialization changes the type of the variable from *void* to the type of the assigned value. Therefore, a single type for a variable in the whole program is insufficient for JavaScript. This requires that the variable types are kept distinct between different positions in the program. This is achieved by using data-flow analysis. The described algorithm combines the Hindley-Milner Algorithm W with abstract interpretation.

4.1 Data Flow Analysis

The control flow graph used for the data-flow analysis is statement-based. For each statement in the program, a control flow graph node is created. Each edge represents a potential control flow between two statements. A node in the control flow graph only represents a basic block if none of the statement's expressions introduce new control flows. For instance, a conditional expression creates two possible control flows — one if the condition is true and another if

Property	Value
Traversal Order	Forward
Node Order	Statement Order
Transfer Function	Hindley-Milner Algorithm W
In- / Out-State	Type Environment Γ
Join Operation	$\Gamma_1 \cup \Gamma_2$
Sensitivity	Flow and Context

Table 1. Properties of the Worklist Algorithm

the condition is not — that are not represented in the control flow graph. The control flow graph is statement-based to reduce the number of states and therefore, the number of states required for the data flow analysis.

The analysis uses the work list algorithm [15] to traverse the control flow nodes in forward order. The analysis is not path-sensitive. The order of the nodes is the same as the order of the statements in the program. The transfer function infers the type for the statement and its expressions using the Hindley-Milner Algorithm W. The in and out state of the data-flow analysis is the type environment. If a node has multiple in branches, then the type environments of these branches are unified. The union of the two type environments contains the mappings of both environments, conflicting mappings are merged using the *unify* function of the Hindley-Milner algorithm. The unification can be defined as follows.

$$\Gamma_1 \cup \Gamma_2 = \{(x, \tau) | x \in \Gamma_1 \vee x \in \Gamma_2\}$$

$$\tau = \begin{cases} \text{unify}(\Gamma_1(x), \Gamma_2(x)) & x \in \Gamma_1 \wedge x \in \Gamma_2 \\ \Gamma_1(x) & x \in \Gamma_1 \wedge x \notin \Gamma_2 \\ \Gamma_2(x) & x \in \Gamma_2 \wedge x \notin \Gamma_1 \end{cases}$$

The important of the worklist algorithm are summarized in the table 1.

4.2 Function Invocation

The algorithm uses inlining for function calls. If a function is called, then the function body is evaluated in the caller's context making the algorithm context-sensitive. Using the type environment of the caller is insufficient for the analysis of the function body as the called function might access variables from its declaration scope. Therefore, the missing mappings from the function declaration type environment Γ_{decl} are added to the caller's type environment Γ_{caller} . This is denoted as $\Gamma_{caller}[\Gamma_{decl}]$. The resulting type environment contains the mappings from both type environments. Mappings already present in the caller's type environment are not overridden with mappings from the function declaration type environment. This is important, as the type of a variable might have changed since the function declaration. This can be defined as follows.

$$\Gamma_{caller}[\Gamma_{decl}] = \{(x, \tau) \mid x \in \Gamma_{caller} \vee x \in \Gamma_{decl}\}$$

$$\tau = \begin{cases} \Gamma_{decl}(x) & x \notin \Gamma_{caller} \\ \Gamma_{caller}(x) & \text{otherwise} \end{cases}$$

4.3 Type System

The type system is designed to infer the types for arbitrary JavaScript code without the need for type annotations. The precision of the inferred type degrades for very dynamic code fragments. The current type system infers the types of all terms and has no support for type annotations. The type system is designed to catch the following errors.

1. Reading of or assigning to an undeclared variable
2. Accessing a property of *null* or *undefined*
3. Invoking a non-function value
4. Invoking a function with missing or incompatible arguments
5. Applying an operator with illegal operands

The defined type system does not distinguish errors by their severity.

Maybe Type The type system distinguishes between absent values, values that may be present, and values that are present for certain. This is needed to detect access to properties on *undefined* or *null* values causing runtime exceptions. The values *null* and *void* are modeled as unit types and represent values that are absent for certain. Potentially absent values are represented by the type *Maybe*<*T*>. The type represents a value that is present for some paths but is not for others. It includes the values *void*, *null*, and all values defined by *T*. Access to a property of a value that may be absent needs to be guarded by a null check.

Record Type Objects supporting members are represented as record types. A record type consists of a set of members. A member is defined by a unique label and the type. The object expression `{ name: "Test", age: 18 }` is represented as record type with two members. The labels of the members are *name* and *age*, the members have the type *string* and *number*.

The type system uses structural typing for record types. Therefore, a type τ_1 is a subtype of τ_2 if τ_1 has at least the same members as τ_2 and the type of each of these members is a subtype of the corresponding member in τ_2 . The type system does not support nominal typing. Nominal typing is required to support classes and prototype based type checking. The type system does not support classes nor prototypes, therefore nominal typing is not supported either.

Array Type The array type *T*[] describes an array containing elements of type *T*. The array type is a specialized record type. The elements contained in an array need to be homogenous. Heterogenous arrays are not supported as the type system does not define a union type. A union type allows

values of different types, e.g., the union “*string* or *number*” contains either *string* or *number* values.

The array type does not track the element type by the element’s position in the array. Accessing element members of an heterogenous array therefore requires a type check if the element is of the expected type. To increase the precision for small, static arrays, a tuple type can be defined. The tuple type tracks the type for every position in the array. Accessing a tuple element therefore does not require an explicit type guard.

Function Type The function type describes a function or method. A function type is characterized by the type of its arguments and its return type, but also the type of the value referenced by *this*. The structure of *this* is defined by the accessed members of *this* inside the function body. The type for *this* is not implicitly defined by the object to which the method belongs.

The presented algorithm uses inlining for invocations of non-native functions. The exact type of a function is not inferred. The inferred type for a function declaration uses type variables for the type of *this*, the arguments, and the return value. Type checking of the function body is performed for every invocation by replacing the type variables with the actual types used in the invocation.

The function type is also used to describe native functions of the host-environment. Type checking an invocation of a native function requires testing if the actual *this* type and the type of the passed arguments are subtypes of the expected type. The *Maybe* allows defining optional arguments for functions.

The implementation does not yet support invoking a function using *call* nor *apply*. Binding the referenced value of *this* using *bind* or the *this* of the outer scope by using arrow functions are not supported either.

Any Type The type system models the special type *any*. The type *any* is a super and subtype of all types. The type inference uses the *any* type as backdoor whenever the type cannot be inferred. Accessing a property of an *any* value yields *any*. An *any* value can be invoked as functions with arbitrary arguments, which returns *any*. Type checking is completely disabled for *any* values. Therefore, an inferred *any* value inherently leads to unsafe code.

Limitations The presented type system supports only a limited set of JavaScript features. It neither has support for classes nor prototype based objects. It does not support the ECMAScript module syntax and therefore all analyzable programs are limited to a single file. Also the type system does not model all features precisely. The plus and compare operators are limited to numbers, even though *strings* can be concatenated using the plus operand and any object with a *valueOf* method can be compared.

5 Evaluation

As part of this work, the tool ESChecker [18] has been implemented, applying the described algorithm. The set of implemented JavaScript features is not sufficient to perform an evaluation using realistic projects in broader practice. The evaluation uses sample listings to compare the implementation with Flow, TAJs, and Infernu¹. Infernu and TAJs do not support ECMAScript 6 and therefore, ECMAScript 5 equivalents of the listings are used.

The test cases and their results are summarized in table 2. Cells marked with a \checkmark indicate that the test case is handled correctly by the type checker specified by the column. An empty cell indicates that the evaluation does not lead to the expected result. The test cases focuses on features that are difficult to analyze statically — as mentioned in section 1 — or features commonly used.

5.1 Variable Redefinement

In JavaScript, values of different types can be assigned to the same variable. Therefore, the same variable can have different types at different positions in the program. The initialization of the variable is deferred in the following listing.

```
1 let x;
2 x = { name: "Micha" };
3
4 x.name;
```

Infernu is the only tool that rejects the program — that is well defined at runtime. The type system used by Infernu is over restrictive and does not allow assigning values to variables where the type of the value is not a sub-type of the variable. In this example, an object value is assigned to a variable of the incompatible type *undefined*.

5.2 Closures

JavaScript supports closures, allowing functions to read and modify variables from the outer scope. The analysis needs to be context-sensitive to support closures.

```
1 let currentId = 0;
2
3 function uniqueId() {
4   return ++currentId;
5 }
6
7 uniqueId();
```

Closures are supported by all evaluated tools.

5.3 Side Effects

Functions in JavaScript do not have to be pure. A function call can have side effects to passed arguments or variables from the enclosing scope of the called function. For type checking, only side affects affecting the type of a variable are relevant. The function called in the following example adds a new property *address* to the object passed as argument. This change of the arguments structure needs to be reflected in the caller's context.

¹The evaluation uses the following versions: Flow 0.24.1, TAJs 0.9-7, Infernu 0.0.0.1.

```
1 function setAddress(p, street, zip) {
2   p.address = { street, zip };
3 }
4
5 const person = { lastName: "Reiser" };
6 setAddress(person, "Bahnhofstrasse 12", 8001);
7 const street = person.address.street;
```

Neither Flow nor Infernu allow side effects changing the structure of an argument and therefore reject the program. Adding new properties is intentionally prohibited by Flow to detect spelling errors.

This example is accepted by TAJs and ESChecker.

5.4 Function Overloading

JavaScript does not provide built-in support for function overloading. A function can only have one definition. But the number of arguments passed to a function must not exactly match the number of declared parameters in the function declaration. This allows optional arguments and therefore function overloading. The following example provides a *range* function that can be called with one or up to three arguments².

```
1 function range(start, end, step) {
2   if (end === undefined) {
3     end = start;
4     start = 0;
5   }
6
7   if (step === undefined) {
8     step = 1;
9   }
10
11   const result = ...;
12   return result;
13 }
14
15 const r = range(10);
16 const r2 = range(1, 10);
17 const r3 = range(10, 1, -1);
```

Infernu does not support optional arguments and therefore rejects the program.

5.5 Callbacks

Callbacks are commonly used in JavaScript for asynchronous and functional programming. Functions can be passed as values. This requires that the flow of function values is tracked. The following example implements the *map* function. The *map* function calls a callback for every array element and puts the result in a new array that is returned. The *map* function is a rank-1 polymorphic function, it can be applied with arguments of arbitrary types. The type inference algorithm needs to infer the principle type and not a specialization [17]. To verify that the principal type is inferred, the *map* function is applied with an array of numbers and a second time with an array of strings.

²The implementation cannot use the default parameters of ES6 as the semantic of the passed arguments depends on the number of arguments. If the function is called with a single argument, then the argument defines the *end* of the range and not the *start*. At the other hand, if *range* is invoked with two arguments, then the first argument specifies the *start* and the second the *end*. Therefore, the semantic of an argument is not only defined by its position but by the position and the number of arguments.

Test Case	Infernu	TAJS	Flow	ESChecker
Variable Redefinement		✓	✓	✓
Closures	✓	✓	✓	✓
Side Effects		✓		✓
Function Overloading		✓	✓	✓
Callbacks	✓	✓		✓
Built-in Types			✓	✓
Dynamic Object Manipulation		✓		
DOM Events		✓		
Frameworks				

Table 2. Evaluation Results

```

1  function map(array, mapper) {
2    const result = [];
3    // ...
4    return result;
5  }
6
7  const array = [1, 2, 3, 4, 5, 6];
8  const doubled = map(array, x => x * 2);
9
10 const names = ["Anisa", "Ardelia", "Madlyn"];
11 const lengths = map(names, name => name.length);

```

Flow does not support type inference for rank-1 polymorphic functions. The example is accepted by Flow if either type annotations are added the *map*, *map* is only applied once, or the passed arrays are of equal types. All other tools support type checking of callbacks and type inference for rank-1 polymorphic functions.

5.6 Built-in Types

JavaScript defines built-in objects and functions. These functions are natively implemented and not available as JavaScript source code, hence the source code cannot be analyzed by the type inference algorithm. This requires that native object and functions are defined in the type checker or externally. The following example uses the native array functions *map* and *reduce*³.

The result of the *reduce* function is an accumulated value. The function uses a callback — passed as first argument — to accumulate the values in the array. The first argument of the callback is the accumulated value over all preceding array elements. The second argument of the callback is the current array element. The callback adds the current element to the accumulated value of the preceding elements and returns the accumulated value. The second argument of *reduce* is the initial accumulator value. The example uses an incorrect order of the arguments passed to the *reduce* function to verify if the type checker validates callbacks passed to built-in functions.

```

1  const array = [1, 2, 3, 4, 5];
2
3  const doubled = array.map(x => x * 2);
4  const sum = array.reduce(0, (accum, x) => accum + x);

```

³As Infernu does not support optional arguments, all callback arguments have been added to the callbacks of *map* and *reduce* before analyzing the example with Infernu.

Infernu rejects the example as it does not support the built-in *reduce* method. TAJs rejects the example as it does not support the built-in *map* method. Flow and ESChecker correctly type the application of the *map* function and reject the application of the *reduce* function.

5.7 Dynamic Object Manipulation

The structure of an object can be dynamically defined or manipulated by using reflection-like code. Such code uses computed property names, something that requires reflection in statically typed languages. The following example is an implementation of the widely used *defaults* function. The *defaults* function accepts two objects. It initializes the properties of the *target* with the values defined in *source*. Properties defined in *source* that are missing in *target* are added. This pattern is commonly used to initialize absent properties with default values for option-objects.

```

1  function defaults(target, source) {
2    target = target === undefined ? {} : target;
3    for (const key of Object.keys(source)) {
4      target[key] = target[key] === undefined ?
5      ↪ source[key] : target[key];
6    }
7    return target;
8  }
9
10 let options = defaults({}, {
11   rounds: 1000,
12   precision: 1
13 });
14 for (let i = 1; i < options.rnds; ++i) {
15   // ...
16 }

```

A precise and sound analysis needs to unroll the for loop on line three to know which properties are copied from the *source* to the *target* object on line four. Unrolling requires that the analysis understands the semantics of *Object.keys*.

The example initializes a plain object with default values on line ten. The initialized option-object is used on line 14, but the property name *rnds* for *rounds* is misspelled. TAJs is the only implementation that detects the misspelled property name if a *for ... in* loop is used instead of the unsupported *Object.keys* method. Flow and ESChecker are unsound for such dynamic code and therefore do not detect the misspelled

property name. Infernu has no support for *for in* loops and therefore fails to type check the program.

5.8 DOM Events

JavaScript is often used in web applications to add dynamic effects by using the DOM-API. A type checker also needs to model the DOM-API, as needed for other built-in types like arrays. Modeling the types and methods defined by the DOM-API is insufficient to achieve type safety, the type checker also needs to model the DOM events [8]. Reacting to user input is done by registering a listener for a particular event on a specific DOM element, like an input field. The API for registering listeners is generic, but the event passed to the listener depends on the type of the handled event. A *keydown* event passes a *KeyboardEvent* object to the listener. The *KeyboardEvent* has additional properties allowing the identification of the pressed key. The type of the event is defined by the event name specified when the listener is registered as shown in the following example on line ten and eleven.

```
1  function onKeyDown(event) {
2    if (event.getModifierState("Shift")) {
3      console.log("Shift...");
4    }
5  }
6
7  const input = document.getElementById("pwd");
8
9  if (input) {
10   input.addEventListener("keydown", onKeyDown,
11     ↳ false);
12   input.addEventListener("blur", onKeyDown, false);
13 }
```

Registering the *onKeyDown* listener for the *keydown* event is legitimate as the *KeyboardEvent* defines the *getModifierState* method used by the listener. On the other hand registering the same listener for the *blur* event is erroneous as the *blur* event does not define the *getModifierState* method.

TAJS is the only tool that detects the malicious invocation of the *getModifierState* method in case of a *blur* event. But TAJS also reports a warning that the variable *input* can potentially be null on line ten and eleven, regardless of the preceding null check on line nine. The same is true for ES-Checker, but it reports an error instead of a warning and therefore rejects the whole program. The reason for this is that ESChecker is neither path-sensitive nor does it evaluate the test condition.

Infernu fails to type check the given example as it does not model the DOM API. Flow does not reject the given program but neither detects the invocation of the not defined function *getModifierState* for the *blur* event.

A special handling for DOM events is beneficial, as it adds additional type safety. But events are not only used when interacting with the DOM, but also if frameworks are being used. Various frameworks like Angular [1] or Ember [3] use events internally and as part of their API. A type checker therefore does not only need to model the DOM events but also the framework events to achieve type safety.

5.9 The Impact of Frameworks on Type Inference

Frameworks are a special challenge for type inference as they make heavy use of dynamic invocations. Frameworks are implemented according to the inversion of control principle, the framework invokes the application, not vica versa. This inverts the control- and data-flow, that are decisive for type inference.

Inferring the correct type is often only possible if the type of the arguments are known. Without an actual invocation, these are unknown. The Angular [1] controller implementation shown in this example loads the persons using the *\$http* service on line ten. The name of the *get* method is misspelled.

```
1  class PersonController {
2    constructor($http) {
3      this.$http = $http;
4      this.persons = [];
5    }
6
7    loadPersons() {
8      this.$http.gt("/persons").then(
9        response => this.persons = response.data);
10   }
11 }
```

The problem shown by this example is that type inference for code using a framework with dynamic invocation is hardly possible for a framework-agnostic type checker as the control and data-flow cannot be inferred using static analysis. The *\$http* service accepted in the constructor on line four is a service that is injected by angular at runtime. The implementation to inject is resolved by name, therefore no explicit connection between the service usage and the service implementation exists. But the type of *\$http* must be known to perform type checking. As the type for *\$http* is not known, its usages cannot be type checked. Explicit type annotations are needed for the *\$http* parameter in the constructor.

5.10 Interpretation of Evaluation Results

The evaluation shows that none of the evaluated tools supports all test cases. ESChecker achieves preciser results than Flow and Infernu, but not as precise as TAJS. TAJS has the most precise type inference and therefore also the most precise type checking results. TAJS preliminary limitation is the lack of support for many ECMAScript 5 features. This might be the major impediment for its adaption for real world projects. Flow at contrary has the best support for ECMAScript 6 script features but type inference is less precise. The precision can be increased by adding type annotations. The type system used by Infernu is over restrictive and rejects most of the test cases, all of them are well typed at runtime. Compared to Infernu, ESChecker has a lower false positive rate and therefore combining the Hindley Milner Algorithm W with data flow analysis has proven to be beneficial for its precision.

But the precision is not sufficient for all examples. Path-sensitivity is needed to access members of a potentially *null* object as the value is potentially null in all branches, independent of their test conditions. Path sensitivity is also needed to support advanced function overloading where a function accepts an argument to be of multiple types, e.g., either a number or a string. The value of the argument is then converted to a common representation in the function body using

type-specific branches. Therefore, the algorithm should be extended to be path sensitive.

The evaluation also showed that the precision for the type inference degrades for very dynamic code, e.g., if the structure of an object is dynamically manipulated. A considerable alternative to type inference for these edge cases — supporting this edge cases might add a high complexity to the type inference algorithm — is the use of type annotations to improve the overall result. Type annotations can either be extracted from JSDoc or be defined in external declaration files, similar to the TypeScript definition files. This increases the precision of the type inference algorithm and therefore the achievable results with type checking. A tool implementing this approach should emit a warning if the type inference algorithm cannot infer the types to encourage the programmer to add type annotations. Allowing any form of type annotations also has the benefit that functions invoked by frameworks can be annotated and therefore can be type checked.

6 Conclusion

The tool support for JavaScript is especially small compared to its popularity. Developers need to relay on manual testing or unit tests for revealing programming errors. The implemented type checker provides a tool that is capable of inferring the types and catch a variety of errors through type checking. The evaluation shows that the analysis is precise and sound for most of the scenarios and sometime provides better results than competitive tools. Therefore, the tool can provide valuable feedback.

But the evaluation also shows that the presented algorithm has its limitations. First, the precision decreases for very dynamic code, for instance, when objects are dynamically manipulated. The precision can be improved if type annotations are used in these cases to substitute type inference. A second limitation is the inability to access properties of potential absent values as the analysis is not path-sensitive and therefore, the value is potentially absent in all branches. Path sensitivity is also required to support type-specific branches, a technique often used to emulate function overloading.

Further, the set of supported features is not sufficient to analyze real-world projects. The implementation is still missing elementary features, such as classes, prototyping or modules. These features are all essential and needed before the tool is useful. Supporting the features defined in ECMAScript 6 and in the upcoming ECMAScript 7 standard requires a tremendous amount of additional work that exceeds the scope of a project thesis by a multitude. But this project thesis shows that precise inference results can be achieved for a majority of JavaScript that, combined with type checking, provides a valuable and immediate feedback to programmers.

References

- [1] *AngularJS — Superheroic JavaScript MVW Framework*. URL: <https://angularjs.org/> (visited on 05/17/2016).
- [2] Anton Ekblad. “Towards a Declarative Web”. MA thesis. University of Gothenburg, 2012. URL: <http://haste-lang.org/pubs/hastereport.pdf> (visited on 04/26/2016).
- [3] *Ember.js*. URL: <http://emberjs.com/> (visited on 05/17/2016).
- [4] Facebook. *Flow, a new static type checker for JavaScript*. Nov. 2014. URL: <https://code.facebook.com/posts/1505962329687926/flow-a-new-static-type-checker-for-javascript/> (visited on 04/26/2016).
- [5] jQuery Foundation. *ESLint*. 2016. URL: <http://eslint.org/> (visited on 04/26/2016).
- [6] Marijn Haverbeke. *Tern.js*. URL: <http://ternjs.net/> (visited on 04/26/2016).
- [7] Ecma International. *ECMAScript® 2015 Language Specification*. Ecma International, June 2015. URL: <http://www.ecma-international.org/ecma-262/6.0/ECMA-262.pdf> (visited on 04/26/2016).
- [8] Simon Holm Jensen, Magnus Madsen, and Anders Møller. “Modeling the HTML DOM and Browser API in Static Analysis of JavaScript Web Applications”. In: *Proc. 8th joint meeting of the European Software Engineering Conference and the ACM SIGSOFT Symposium on the Foundations of Software Engineering (ESEC/FSE)*. Sept. 2011. URL: <http://cs.au.dk/~amoeller/papers/dom/paper.pdf> (visited on 04/26/2016).
- [9] Simon Holm Jensen, Anders Møller, and Peter Thiemann. “Type Analysis for JavaScript”. In: *Proc. 16th International Static Analysis Symposium (SAS)*. Vol. 5673. LNCS. Springer-Verlag, Aug. 2009. URL: <http://cs.au.dk/~amoeller/papers/tajs/paper.pdf> (visited on 04/26/2016).
- [10] *JSDoc*. URL: <http://usejsdoc.org/> (visited on 05/22/2016).
- [11] Noam Lewis. *Inernu*. URL: <https://noamlewis.wordpress.com/2015/01/20/introducing-sjs-a-type-inferer-and-checker-for-javascript/> (visited on 04/26/2016).
- [12] Brian McKenna. *Roy*. URL: <http://roy.brianmckenna.org/> (visited on 04/26/2016).
- [13] Microsoft. *TypeScript*. 2012. URL: <https://www.typescriptlang.org/> (visited on 04/26/2016).
- [14] Robin Milner. “A theory of type polymorphism in programming”. In: *Journal of Computer and System Sciences* 17 (1978), pp. 348–375.
- [15] Flemming Nielson, Hanne R. Nielson, and Chris Hankin. *Principles of Program Analysis*. Secaucus, NJ, USA: Springer-Verlag New York, Inc., 1999. ISBN: 3540654100.
- [16] Morten Passow Odgaard. “JavaScript Type Inference Using Dynamic Analysis”. MA thesis. Aarhus University, 2014. URL: http://cs.au.dk/fileadmin/site_files/cs/Masters_

and _diplomas / MortenPassowOdgaard .
pdf (visited on 04/26/2016).

- [17] Benjamin C. Pierce. *Types and Programming Languages*. 1st. The MIT Press, 2002. ISBN: 0262162091, 9780262162098.
- [18] Micha Reiser. *ESChecker*. online. 2016. URL: <https://github.com/DatenMetzgerX/ESChecker>.